

Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
import warnings
warnings.simplefilter('ignore')
```

Data Collection

```
In [2]: # Loading the dataset to a Pandas DataFrame
wine_data = pd.read_csv(r'C:\Users\ADMIN\Desktop\Projects\Data sets\winequality-red.csv')
```

```
In [5]: wine_data
```

```
Out[5]:
```

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol | quality |
|------|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|------|-----------|---------|---------|
| 0 | 7.4 | 0.700 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.99780 | 3.51 | 0.56 | 9.4 | 5 |
| 1 | 7.8 | 0.880 | 0.00 | 2.6 | 0.098 | 25.0 | 67.0 | 0.99680 | 3.20 | 0.68 | 9.8 | 5 |
| 2 | 7.8 | 0.760 | 0.04 | 2.3 | 0.092 | 15.0 | 54.0 | 0.99700 | 3.26 | 0.65 | 9.8 | 5 |
| 3 | 11.2 | 0.280 | 0.56 | 1.9 | 0.075 | 17.0 | 60.0 | 0.99800 | 3.16 | 0.58 | 9.8 | 6 |
| 4 | 7.4 | 0.700 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.99780 | 3.51 | 0.56 | 9.4 | 5 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1594 | 6.2 | 0.600 | 0.08 | 2.0 | 0.090 | 32.0 | 44.0 | 0.99490 | 3.45 | 0.58 | 10.5 | 5 |
| 1595 | 5.9 | 0.550 | 0.10 | 2.2 | 0.062 | 39.0 | 51.0 | 0.99512 | 3.52 | 0.76 | 11.2 | 6 |
| 1596 | 6.3 | 0.510 | 0.13 | 2.3 | 0.076 | 29.0 | 40.0 | 0.99574 | 3.42 | 0.75 | 11.0 | 6 |
| 1597 | 5.9 | 0.645 | 0.12 | 2.0 | 0.075 | 32.0 | 44.0 | 0.99547 | 3.57 | 0.71 | 10.2 | 5 |
| 1598 | 6.0 | 0.310 | 0.47 | 3.6 | 0.067 | 18.0 | 42.0 | 0.99549 | 3.39 | 0.66 | 11.0 | 6 |

1599 rows × 12 columns

```
In [3]: wine_data.shape
```

```
Out[3]: (1599, 12)
```

```
In [4]: wine_data.head()
```

```
Out[4]:
```

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol | quality |
|---|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|------|-----------|---------|---------|
| 0 | 7.4 | 0.70 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 | 9.4 | 5 |
| 1 | 7.8 | 0.88 | 0.00 | 2.6 | 0.098 | 25.0 | 67.0 | 0.9968 | 3.20 | 0.68 | 9.8 | 5 |
| 2 | 7.8 | 0.76 | 0.04 | 2.3 | 0.092 | 15.0 | 54.0 | 0.9970 | 3.26 | 0.65 | 9.8 | 5 |
| 3 | 11.2 | 0.28 | 0.56 | 1.9 | 0.075 | 17.0 | 60.0 | 0.9980 | 3.16 | 0.58 | 9.8 | 6 |
| 4 | 7.4 | 0.70 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 | 9.4 | 5 |

Checking missing Values

```
In [6]: wine_data.isnull().sum()
```

```
Out[6]: fixed acidity      0
volatile acidity    0
```

```
citric acid      0
residual sugar   0
chlorides        0
free sulfur dioxide  0
total sulfur dioxide  0
density          0
pH              0
sulphates        0
alcohol          0
quality          0
dtype: int64
```

Statistical measures


```
In [7]: wine_data.describe()
```

Out[7]:

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | |
|-------|---------------|------------------|-------------|----------------|-------------|---------------------|----------------------|-------------|-------------|-------------|-------------|
| count | 1599.000000 | 1599.000000 | 1599.000000 | 1599.000000 | 1599.000000 | 1599.000000 | 1599.000000 | 1599.000000 | 1599.000000 | 1599.000000 | 1599.000000 |
| mean | 8.319637 | 0.527821 | 0.270976 | 2.538806 | 0.087467 | 15.874922 | 46.467792 | 0.996747 | 3.311113 | 0.658149 | 1599.000000 |
| std | 1.741096 | 0.179060 | 0.194801 | 1.409928 | 0.047065 | 10.460157 | 32.895324 | 0.001887 | 0.154386 | 0.169507 | 1599.000000 |
| min | 4.600000 | 0.120000 | 0.000000 | 0.900000 | 0.012000 | 1.000000 | 6.000000 | 0.990070 | 2.740000 | 0.330000 | 1599.000000 |
| 25% | 7.100000 | 0.390000 | 0.090000 | 1.900000 | 0.070000 | 7.000000 | 22.000000 | 0.995600 | 3.210000 | 0.550000 | 1599.000000 |
| 50% | 7.900000 | 0.520000 | 0.260000 | 2.200000 | 0.079000 | 14.000000 | 38.000000 | 0.996750 | 3.310000 | 0.620000 | 1599.000000 |
| 75% | 9.200000 | 0.640000 | 0.420000 | 2.600000 | 0.090000 | 21.000000 | 62.000000 | 0.997835 | 3.400000 | 0.730000 | 1599.000000 |
| max | 15.900000 | 1.580000 | 1.000000 | 15.500000 | 0.611000 | 72.000000 | 289.000000 | 1.003690 | 4.010000 | 2.000000 | 1599.000000 |


```
In [8]: sns.catplot(x='quality',data = wine_data, kind = 'count')
```


Out[8]: <seaborn.axisgrid.FacetGrid at 0x219e25c3af0>



```
In [9]: # volatile acidity vs quality
plot = plt.figure(figsize=(5,5))
sns.barplot(x='quality',y='volatile acidity',data=wine_data)
```


Out[9]: <AxesSubplot:xlabel='quality', ylabel='volatile acidity'>





```
In [10]: # citric acid vs quality
plot = plt.figure(figsize=(5,5))
sns.barplot(x='quality',y='citric acid',data=wine_data)
```

```
Out[10]: <AxesSubplot:xlabel='quality', ylabel='citric acid'>
```




Correlation


1. Positive Correlation
2. Negative Correlation

```
In [12]: correlation = wine_data.corr()
```

```
In [13]: # constructing a heatmap to understand the correlation between the columns
plt.figure(figsize=(10,10))
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f', annot=True, annot_kws={'size':8}, cmap='Blues')
```

```
Out[13]: <AxesSubplot:>
```





Data PreProcessing

```
In [14]: X = wine_data.drop('quality',axis=1)
```

```
In [16]: print(X)
```

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | \ |
|------|---------------|------------------|-------------|----------------|-----------|---|
| 0 | 7.4 | 0.700 | 0.00 | 1.9 | 0.076 | |
| 1 | 7.8 | 0.880 | 0.00 | 2.6 | 0.098 | |
| 2 | 7.8 | 0.760 | 0.04 | 2.3 | 0.092 | |
| 3 | 11.2 | 0.280 | 0.56 | 1.9 | 0.075 | |
| 4 | 7.4 | 0.700 | 0.00 | 1.9 | 0.076 | |
| ... | ... | ... | ... | ... | ... | |
| 1594 | 6.2 | 0.600 | 0.08 | 2.0 | 0.090 | |
| 1595 | 5.9 | 0.550 | 0.10 | 2.2 | 0.062 | |
| 1596 | 6.3 | 0.510 | 0.13 | 2.3 | 0.076 | |
| 1597 | 5.9 | 0.645 | 0.12 | 2.0 | 0.075 | |
| 1598 | 6.0 | 0.310 | 0.47 | 3.6 | 0.067 | |

| | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | \ |
|------|---------------------|----------------------|---------|------|-----------|---|
| 0 | 11.0 | 34.0 | 0.99780 | 3.51 | 0.56 | |
| 1 | 25.0 | 67.0 | 0.99680 | 3.20 | 0.68 | |
| 2 | 15.0 | 54.0 | 0.99700 | 3.26 | 0.65 | |
| 3 | 17.0 | 60.0 | 0.99800 | 3.16 | 0.58 | |
| 4 | 11.0 | 34.0 | 0.99780 | 3.51 | 0.56 | |
| ... | ... | ... | ... | ... | ... | |
| 1594 | 32.0 | 44.0 | 0.99490 | 3.45 | 0.58 | |
| 1595 | 39.0 | 51.0 | 0.99512 | 3.52 | 0.76 | |
| 1596 | 29.0 | 40.0 | 0.99574 | 3.42 | 0.75 | |
| 1597 | 32.0 | 44.0 | 0.99547 | 3.57 | 0.71 | |
| 1598 | 18.0 | 42.0 | 0.99549 | 3.39 | 0.66 | |

| | alcohol |
|------|---------|
| 0 | 9.4 |
| 1 | 9.8 |
| 2 | 9.8 |
| 3 | 9.8 |
| 4 | 9.4 |
| ... | ... |
| 1594 | 10.5 |
| 1595 | 11.2 |
| 1596 | 11.0 |
| 1597 | 10.2 |
| 1598 | 11.0 |

[1599 rows x 11 columns]

Label Binarization

```
In [17]: Y = wine_data['quality'].apply(lambda y_value: 1 if y_value >= 7 else 0)
```

```
In [18]: print (Y)
```

| | |
|------|-----|
| 0 | 0 |
| 1 | 0 |
| 2 | 0 |
| 3 | 0 |
| 4 | 0 |
| ... | ... |
| 1594 | 0 |
| 1595 | 0 |
| 1596 | 0 |
| 1597 | 0 |

1598 0
Name: quality, Length: 1599, dtype: int64

Train and Test Split

```
In [19]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=3)
```

```
In [20]: print(Y.shape, Y_train.shape, Y_test.shape)

(1599,) (1279,) (320,)
```

Model Training:

Random Forest Classifier

```
In [21]: model = RandomForestClassifier()
```

```
In [22]: model.fit(X_train, Y_train)
```

```
Out[22]: RandomForestClassifier()
```

Model Evaluation

Accuracy score

```
In [23]: # Accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

```
In [24]: print('Accuracy:', test_data_accuracy)
```

Accuracy: 0.928125

Building a Predictive System

```
In [25]: input_data1 = (7.5,0.5,0.36,6.1,0.071,17.0,102.0,0.9978,3.35,0.8,10.5)

# Changing the input data in to a numpy array

input_data_as_numpy_array = np.asarray(input_data1)

# Reshape the data as we are predicting the label for only one instance
input_data_resaped = input_data_as_numpy_array.reshape(1,-1)

prediction = model.predict(input_data_resaped)
print(prediction)

if prediction[0]==1:
    print('Good Quality Wine')
else:
    print('Bad Quality Wine')
```

[0]
Bad Quality Wine

```
In [26]: input_data = (7.3,0.65,0.0,1.2,0.065,15.0,21.0,0.9946,3.39,0.47,10.0)
```

```
# Changing the input data in to a numpy array
input_data_as_numpy_array = np.asarray(input_data)

# Reshape the data as we are predicting the label for only one instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

prediction = model.predict(input_data_reshaped)
print(prediction)

if prediction[0]==1:
    print('Good Quality Wine')
else:
    print('Bad Quality Wine')
```

```
[1]
Good Quality Wine
```

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js